

# Alzheimer\_Prediction\_with\_Demographics\_Data\_\_\_\_\_

November 4, 2021

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sn
from sklearn import preprocessing
from sklearn.metrics import log_loss
from sklearn.metrics import confusion_matrix
import sklearn.metrics as metrics
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import classification_report
%matplotlib inline
```

Load The data files    Display data frame head

```
[2]: data_frame = pd.read_excel('C:/Users/amran/Downloads/
↳oasis_longitudinal_demographics.xlsx')

data_frame.head()
```

```
[2]:  Subject ID      MRI ID      Group  Visit  MR Delay  M/F  Hand  Age  EDUC  \
0  OAS2_0001  OAS2_0001_MR1  Nondemented      1         0    M    R   87   14
1  OAS2_0001  OAS2_0001_MR2  Nondemented      2        457    M    R   88   14
2  OAS2_0002  OAS2_0002_MR1    Demented      1         0    M    R   75   12
3  OAS2_0002  OAS2_0002_MR2    Demented      2        560    M    R   76   12
4  OAS2_0002  OAS2_0002_MR3    Demented      3       1895    M    R   80   12

      SES  MMSE  CDR      eTIV      nWBV      ASF
0    2.0   27.0   0.0  1986.550000  0.696106  0.883440
1    2.0   30.0   0.0  2004.479526  0.681062  0.875539
2   NaN   23.0   0.5  1678.290000  0.736336  1.045710
3   NaN   28.0   0.5  1737.620000  0.713402  1.010000
4   NaN   22.0   0.5  1697.911134  0.701236  1.033623
```

Data pre-processing and selection    Lets first look at columns data types

```
[3]: data_frame.dtypes
```

```
[3]: Subject ID      object
      MRI ID        object
      Group         object
      Visit         int64
      MR Delay      int64
      M/F           object
      Hand          object
      Age           int64
      EDUC          int64
      SES           float64
      MMSE          float64
      CDR           float64
      eTIV          float64
      nWBV          float64
      ASF           float64
      dtype: object
```

```
[4]: data_frame.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 373 entries, 0 to 372
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Subject ID  373 non-null   object
1   MRI ID      373 non-null   object
2   Group       373 non-null   object
3   Visit       373 non-null   int64
4   MR Delay    373 non-null   int64
5   M/F         373 non-null   object
6   Hand        373 non-null   object
7   Age         373 non-null   int64
8   EDUC        373 non-null   int64
9   SES         354 non-null   float64
10  MMSE        371 non-null   float64
11  CDR         373 non-null   float64
12  eTIV        373 non-null   float64
13  nWBV        373 non-null   float64
14  ASF         373 non-null   float64
dtypes: float64(6), int64(4), object(5)
memory usage: 43.8+ KB
```

### Histogram

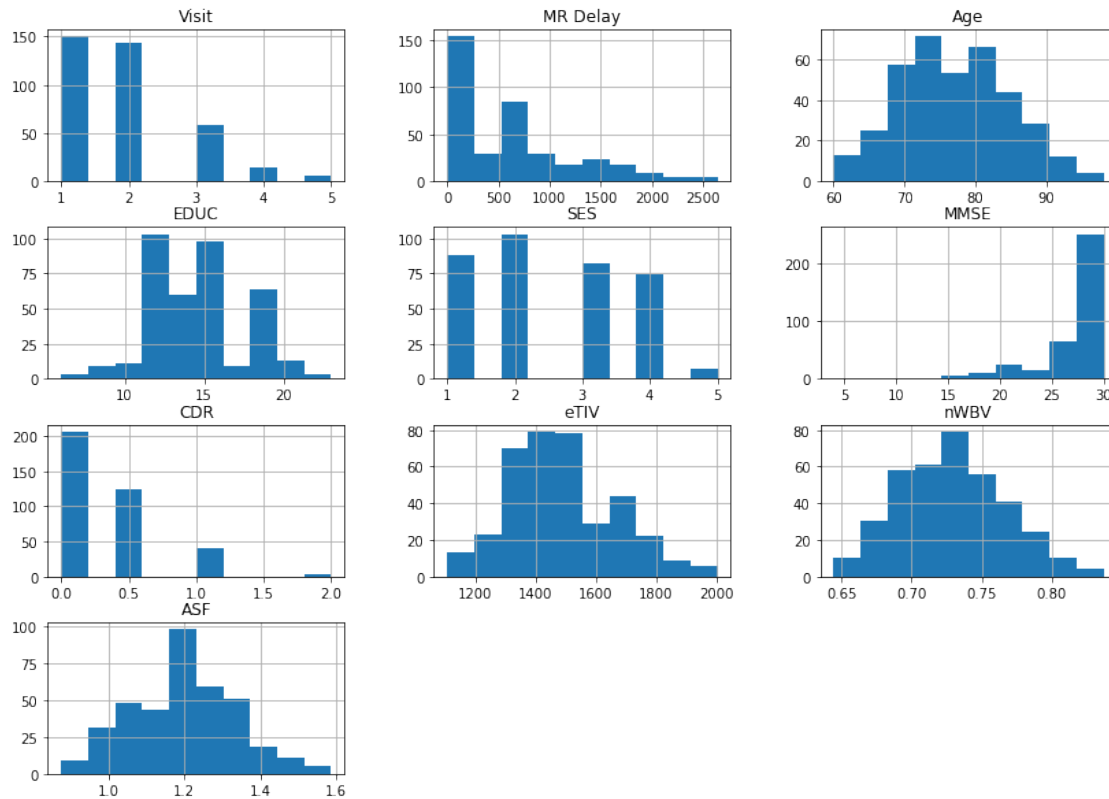
```
[5]: data_frame.hist(figsize =(14,10))
```

```
[5]: array([[<AxesSubplot:title={'center':'Visit'}>,
          <AxesSubplot:title={'center':'MR Delay'}>],
```

```

<AxesSubplot:title={'center':'Age'}>],
[<AxesSubplot:title={'center':'EDUC'}>,
<AxesSubplot:title={'center':'SES'}>,
<AxesSubplot:title={'center':'MMSE'}>],
[<AxesSubplot:title={'center':'CDR'}>,
<AxesSubplot:title={'center':'eTIV'}>,
<AxesSubplot:title={'center':'nWBV'}>],
[<AxesSubplot:title={'center':'ASF'}>, <AxesSubplot:>,
<AxesSubplot:>]], dtype=object)

```



### Correlation Heatmap

```

[6]: import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize = (14,10))


sns.heatmap(data_frame[data_frame.columns].corr(), annot=True)

```

```

[6]: <AxesSubplot:>

```



### Dealing the missing values

```
[7]: missing_data = data_frame.isnull()
missing_data.head()
```

```
[7]:
```

	Subject ID	MRI ID	Group	Visit	MR Delay	M/F	Hand	Age	EDUC	\
0	False	False	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	False	False	

	SES	MMSE	CDR	eTIV	nWBV	ASF
0	False	False	False	False	False	False
1	False	False	False	False	False	False
2	True	False	False	False	False	False
3	True	False	False	False	False	False
4	True	False	False	False	False	False

Count missing values in each column

```
[8]: for column in missing_data.columns.values.tolist():  
      print(column)  
      print(missing_data[column].value_counts())  
      print(" ")
```

Subject ID

False 373

Name: Subject ID, dtype: int64

MRI ID

False 373

Name: MRI ID, dtype: int64

Group

False 373

Name: Group, dtype: int64

Visit

False 373

Name: Visit, dtype: int64

MR Delay

False 373

Name: MR Delay, dtype: int64

M/F

False 373

Name: M/F, dtype: int64

Hand

False 373

Name: Hand, dtype: int64

Age

False 373

Name: Age, dtype: int64

EDUC

False 373

Name: EDUC, dtype: int64

SES

False 354

True 19

Name: SES, dtype: int64

MMSE

False 371

```
True      2
Name: MMSE, dtype: int64
```

```
CDR
False     373
Name: CDR, dtype: int64
```

```
eTIV
False     373
Name: eTIV, dtype: int64
```

```
nWBV
False     373
Name: nWBV, dtype: int64
```

```
ASF
False     373
Name: ASF, dtype: int64
```

SES have 19 missing value and MMSE have 2 missing Value replace missing values

```
[9]: # check the details value of SES
data_frame['SES'].value_counts()
```

```
[9]: 2.0      103
     1.0       88
     3.0       82
     4.0       74
     5.0        7
     Name: SES, dtype: int64
```

```
[10]: #Here 2 is most common values. use the ".idxmax()" method to calculate for us
      ↳ the most common type automatically:
```

```
data_frame['SES'].value_counts().idxmax()
```

```
[10]: 2.0
```

```
[11]: #replace null with most common values
data_frame['SES'].fillna(2.0, inplace=True)
```

```
[12]: #check the details value of MMSE
data_frame['MMSE'].value_counts()
```

```
[12]: 30.0      114
     29.0       91
     28.0       45
```

```

27.0    32
26.0    20
25.0    12
23.0    11
21.0    11
20.0     7
22.0     7
17.0     5
24.0     4
16.0     3
19.0     3
18.0     2
15.0     2
 7.0     1
 4.0     1
Name: MMSE, dtype: int64

```

```

[13]: #30 is the most common value. We can also use the ".idxmax()" method to
      ↪ calculate for us the most common type automatically:
      data_frame['MMSE'].value_counts().idxmax()

```

```

[13]: 30.0

```

```

[14]: #replace null with most common values
      data_frame['MMSE'].fillna(30,inplace=True)

```

```

[15]: missing_data = data_frame.isnull()
      missing_data.head()

```

```

[15]:   Subject ID  MRI ID  Group  Visit  MR Delay  M/F  Hand  Age  EDUC  \
0      False   False  False  False    False  False  False  False  False
1      False   False  False  False    False  False  False  False  False
2      False   False  False  False    False  False  False  False  False
3      False   False  False  False    False  False  False  False  False
4      False   False  False  False    False  False  False  False  False

      SES  MMSE  CDR  eTIV  nWBV  ASF
0  False  False  False  False  False  False
1  False  False  False  False  False  False
2  False  False  False  False  False  False
3  False  False  False  False  False  False
4  False  False  False  False  False  False

```

```

[16]: missing_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 373 entries, 0 to 372
Data columns (total 15 columns):

```

#	Column	Non-Null Count	Dtype
0	Subject ID	373 non-null	bool
1	MRI ID	373 non-null	bool
2	Group	373 non-null	bool
3	Visit	373 non-null	bool
4	MR Delay	373 non-null	bool
5	M/F	373 non-null	bool
6	Hand	373 non-null	bool
7	Age	373 non-null	bool
8	EDUC	373 non-null	bool
9	SES	373 non-null	bool
10	MMSE	373 non-null	bool
11	CDR	373 non-null	bool
12	eTIV	373 non-null	bool
13	nWBV	373 non-null	bool
14	ASF	373 non-null	bool

dtypes: bool(15)

memory usage: 5.6 KB

```
[17]: for column in missing_data.columns.values.tolist():
        print(column)
        print(missing_data[column].value_counts())
        print(" ")
```

Subject ID

False 373

Name: Subject ID, dtype: int64

MRI ID

False 373

Name: MRI ID, dtype: int64

Group

False 373

Name: Group, dtype: int64

Visit

False 373

Name: Visit, dtype: int64

MR Delay

False 373

Name: MR Delay, dtype: int64

M/F

False 373

Name: M/F, dtype: int64



```
Hand
False      373
Name: Hand, dtype: int64
```

```
Age
False      373
Name: Age, dtype: int64
```

```
EDUC
False      373
Name: EDUC, dtype: int64
```

```
SES
False      373
Name: SES, dtype: int64
```

```
MMSE
False      373
Name: MMSE, dtype: int64
```

```
CDR
False      373
Name: CDR, dtype: int64
```

```
eTIV
False      373
Name: eTIV, dtype: int64
```

```
nWBV
False      373
Name: nWBV, dtype: int64
```

```
ASF
False      373
Name: ASF, dtype: int64
```

Let's see how many of each class is in our data set

```
[18]: data_frame['Subject ID'].value_counts()
```

```
[18]: OAS2_0073      5
      OAS2_0048      5
      OAS2_0127      5
      OAS2_0070      5
      OAS2_0037      4
      ..
```

```

OAS2_0128    2
OAS2_0120    2
OAS2_0112    2
OAS2_0086    2
OAS2_0063    2
Name: Subject ID, Length: 150, dtype: int64

```

```
[19]: # check the total number of subject
data_frame['Subject ID'].nunique()
```

```
[19]: 150
```

```
[20]: data_frame['Group'].value_counts()
```

```

[20]: Nondemented    190
      Demented      146
      Converted      37
      Name: Group, dtype: int64

```

#### Check for male and female

```
[21]: data_frame.groupby(['M/F', 'Group'])['Subject ID'].nunique()
```

```

[21]: M/F  Group
      F    Converted    10
           Demented    28
           Nondemented  50
      M    Converted     4
           Demented    36
           Nondemented  22
      Name: Subject ID, dtype: int64

```

#### Group by Male female and Subject ID

```
[22]: data_frame_sub =data_frame.groupby(['M/F', 'Group'])
```

```
[23]: data_frame_sub= data_frame_sub.agg({'Subject ID': 'nunique'})
```

```
[24]: data_frame_sub.head()
```

```

[24]:           Subject ID
      M/F Group
      F    Converted    10
           Demented    28
           Nondemented  50
      M    Converted     4
           Demented    36

```

```
[25]: # Change M to 1 and F to 0
data_frame['M/F'] = data_frame['M/F'].replace(['F','M'], [0,1])

[26]: # Hand cloumn value to numaric data R to 1 and L to 0
data_frame['Hand'] = data_frame['Hand'].replace(['L','R'], [0,1]) # Hand column

[27]: # Group :Replace Converted to Demented
data_frame['Group'] = data_frame['Group'].replace(['Converted'], ['Demented'])

[28]: data_frame.head()
```

```
[28]:
```

	Subject ID	MRI ID	Group	Visit	MR Delay	M/F	Hand	Age	\
0	OAS2_0001	OAS2_0001_MR1	Nondemented	1	0	1	1	87	
1	OAS2_0001	OAS2_0001_MR2	Nondemented	2	457	1	1	88	
2	OAS2_0002	OAS2_0002_MR1	Demented	1	0	1	1	75	
3	OAS2_0002	OAS2_0002_MR2	Demented	2	560	1	1	76	
4	OAS2_0002	OAS2_0002_MR3	Demented	3	1895	1	1	80	

	EDUC	SES	MMSE	CDR	eTIV	nWBV	ASF
0	14	2.0	27.0	0.0	1986.550000	0.696106	0.883440
1	14	2.0	30.0	0.0	2004.479526	0.681062	0.875539
2	12	2.0	23.0	0.5	1678.290000	0.736336	1.045710
3	12	2.0	28.0	0.5	1737.620000	0.713402	1.010000
4	12	2.0	22.0	0.5	1697.911134	0.701236	1.033623

```
[29]: data_frame.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 373 entries, 0 to 372
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Subject ID      373 non-null    object
1   MRI ID          373 non-null    object
2   Group           373 non-null    object
3   Visit           373 non-null    int64
4   MR Delay        373 non-null    int64
5   M/F             373 non-null    int64
6   Hand            373 non-null    int64
7   Age             373 non-null    int64
8   EDUC            373 non-null    int64
9   SES             373 non-null    float64
10  MMSE            373 non-null    float64
11  CDR             373 non-null    float64
12  eTIV            373 non-null    float64
13  nWBV            373 non-null    float64
14  ASF             373 non-null    float64
dtypes: float64(6), int64(6), object(3)
```

memory usage: 43.8+ KB

New dataframe for male female group with Unique Subject ID

```
[30]: data_frame_new = data_frame.groupby(['M/F', 'Group'])
```

```
[31]: data_frame_new = data_frame_new.agg({'Subject ID': 'nunique'})
```

```
[32]: data_frame_new.head()
```

```
[32]:
```

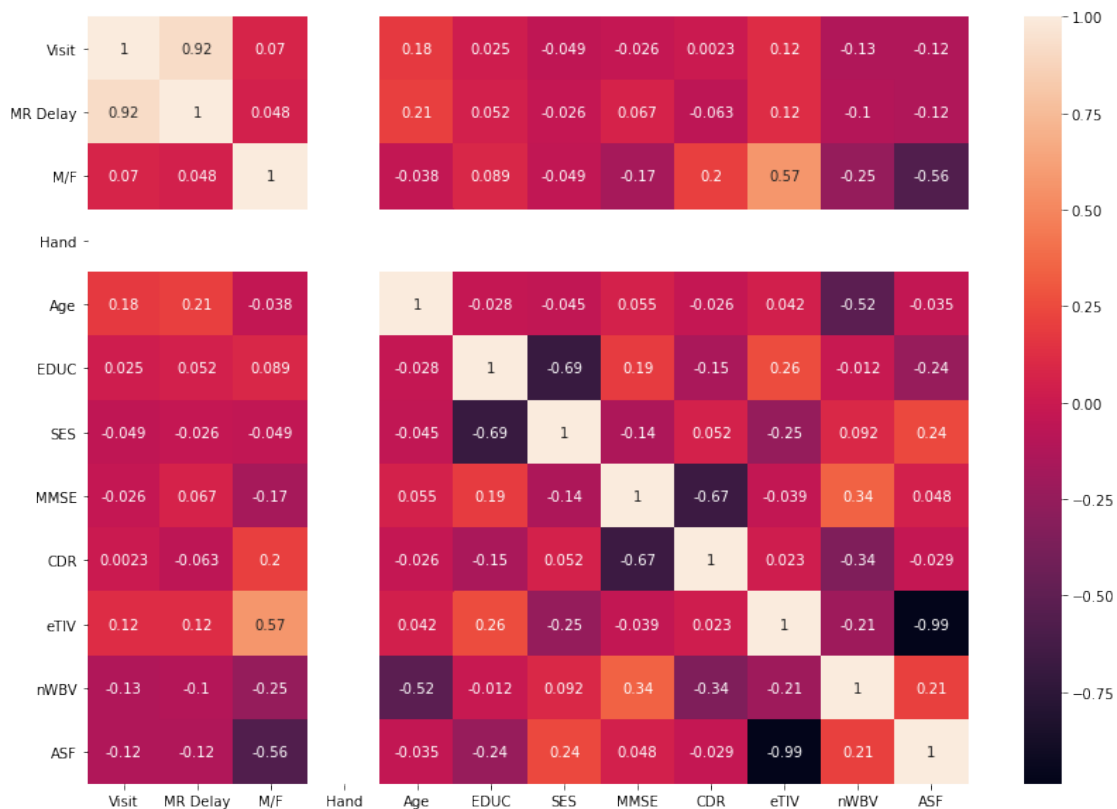
		Subject ID
M/F Group		
0	Demented	38
	Nondemented	50
1	Demented	40
	Nondemented	22

```
[33]: import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize = (14,10))

sns.heatmap(data_frame[data_frame.columns].corr(), annot=True)
```

```
[33]: <AxesSubplot:>
```



Target value, response variable or dependent variable

```
[34]: y_data = data_frame['Group']
```

```
[35]: y_data[0:5]
```

```
[35]: 0    Nondemented
      1    Nondemented
      2      Demented
      3      Demented
      4      Demented
      Name: Group, dtype: object
```

```
[36]: # Independent variable or regressor as X
      # drop unrelated values
      X_data = data_frame.drop(['Subject ID', 'MRI ID', 'Group'], axis=1)
      X_data.head()
```

```
[36]:   Visit  MR Delay  M/F  Hand  Age  EDUC  SES  MMSE  CDR      eTIV  \
0      1         0    1     1   87   14  2.0  27.0  0.0  1986.550000
1      2       457    1     1   88   14  2.0  30.0  0.0  2004.479526
2      1         0    1     1   75   12  2.0  23.0  0.5  1678.290000
3      2       560    1     1   76   12  2.0  28.0  0.5  1737.620000
4      3      1895    1     1   80   12  2.0  22.0  0.5  1697.911134

      nWBV      ASF
0  0.696106  0.883440
1  0.681062  0.875539
2  0.736336  1.045710
3  0.713402  1.010000
4  0.701236  1.033623
```

### 0.0.1 Train and Test Split

```
[37]: # import train_test_split library
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import accuracy_score
      from sklearn.metrics import roc_auc_score
      from sklearn.metrics import roc_curve, auc

      X_train, X_test, y_train, y_test = train_test_split(X_data, y_data, test_size = 0.33)
      print ('Train set:', X_train.shape, y_train.shape)
      print ('Test set:', X_test.shape, y_test.shape)
```

Train set: (249, 12) (249,)  
Test set: (124, 12) (124,)

**Normalize Data** Data Standardization give data zero mean and unit variance

```
[38]: X_data= preprocessing.StandardScaler().fit(X_data).transform(X_data)
      X_data[0:5]
```

```
[38]: array([[ -0.95706686, -0.93771494,  1.15379808,  0.          ,  1.30873772,
        -0.20813199, -0.3944662 , -0.09706416, -0.77765291,  2.83359462,
        -0.90181966, -2.26232493],
       [ 0.12799678, -0.21761337,  1.15379808,  0.          ,  1.43978716,
        -0.20813199, -0.3944662 ,  0.71958842, -0.77765291,  2.93552502,
        -1.30741435, -2.31961167],
       [-0.95706686, -0.93771494,  1.15379808,  0.          , -0.26385558,
        -0.90439416, -0.3944662 , -1.18593426,  0.55905002,  1.08111854,
         0.18280374, -1.08577527],
       [ 0.12799678, -0.05531476,  1.15379808,  0.          , -0.13280614,
        -0.90439416, -0.3944662 ,  0.17515337,  0.55905002,  1.41841305,
        -0.43550978, -1.34469305],
       [ 1.21306043,  2.04826424,  1.15379808,  0.          ,  0.39139163,
        -0.90439416, -0.3944662 , -1.45815179,  0.55905002,  1.19266583,
        -0.76351198, -1.17341289]])
```

## 0.0.2 Support Vector Machine (SVM)

```
[39]: # import SVM library
```

```
      from sklearn import svm
```

```
[40]: svm_clf = svm.SVC(C=1, kernel='linear').fit(X_train, y_train)
```


```
[41]: yhat_svm = svm_clf.predict(X_test)
```

```
[42]: SVM_accuracy_score = accuracy_score(y_test,yhat_svm)*100
      print("SVM_accuracy_score:", SVM_accuracy_score)
```

SVM\_accuracy\_score: 96.7741935483871

### SVM Confusion matrix

```
[43]: #confusion_matrix(y_test,yhat_svm)
      from sklearn.metrics import confusion_matrix
      plot_confusion_matrix(svm_clf, X_test, y_test)
      plt.show()
```



```
[44]: # Confusion matrix using crosstab method of pandas.
svm_pd = pd.crosstab(y_test, yhat_svm, rownames=['True'],
    ↳ colnames=['Predicted'], margins=True)
svm_pd
```

```
[44]: Predicted    Demented  Nondemented  All
True
Demented         53           3      56
Nondemented       1          67      68
All              54          70     124
```

### SVM Classification Report

```
[45]: #print(classification_report(y_test,yhat_svm))

print("confusion matrix\n",confusion_matrix(y_test,yhat_svm))
print("\nclassification report\n",classification_report(y_test,yhat_svm))
print("Accuracy: ",accuracy_score(y_test, yhat_svm)*100)
```

```
confusion matrix
[[53  3]
 [ 1 67]]
```

```
classification report
           precision    recall  f1-score   support
```


Demented	0.98	0.95	0.96	56
Nondemented	0.96	0.99	0.97	68
accuracy			0.97	124
macro avg	0.97	0.97	0.97	124
weighted avg	0.97	0.97	0.97	124

Accuracy: 96.7741935483871

```
[46]: SVM_roc_auc_score = roc_auc_score(y_test, svm_clf.decision_function(X_test),
    ↪average=None)
print("SVM_roc_auc_score:", SVM_roc_auc_score)
```

SVM\_roc\_auc\_score: 0.9684873949579832

```
[47]: metrics.plot_roc_curve(svm_clf, X_test, y_test)
plt.show()
```



### 0.0.3 Logistic Regression

```
[48]: #import Library LogisticRegression
from sklearn.linear_model import LogisticRegression
```



```
[49]: #Train
lr_clf = LogisticRegression(C=1, solver='liblinear').fit(X_train,y_train)
lr_clf
```

```
[49]: LogisticRegression(C=1, solver='liblinear')
```

```
[50]: yhat_lr = lr_clf.predict(X_test)

lr_yhat_prob = lr_clf.predict_proba(X_test)
```


```
[51]: LR_accuracy_score = accuracy_score(y_test,yhat_lr)*100
print("LR_accuracy_score :",LR_accuracy_score )
```

```
LR_accuracy_score : 96.7741935483871
```

### LR Confusion matrix

```
[52]: #confusion_matrix(y_test,yhat_lr)

plot_confusion_matrix(lr_clf, X_test, y_test)
plt.show()
```



```
[53]: # Confusion matrix using crosstab method of pandas.
LR_pd = pd.crosstab(y_test, yhat_lr, rownames=['True'], colnames=['Predicted'],
↪ margins=True)
LR_pd
```

```
[53]: Predicted      Demented  Nondemented  All
      True
      Demented      53           3      56
      Nondemented    1          67      68
      All            54          70     124
```

### LR Classification Report

```
[54]: #print(classification_report(y_test,yhat_lr))

print("confusion matrix\n",confusion_matrix(y_test,yhat_lr))
print("\nclassification report\n",classification_report(y_test,yhat_lr))
print("Accuracy: ",accuracy_score(y_test, yhat_lr)*100)
```

confusion matrix

```
[[53  3]
 [ 1 67]]
```

classification report


	precision	recall	f1-score	support
Demented	0.98	0.95	0.96	56
Nondemented	0.96	0.99	0.97	68
accuracy			0.97	124
macro avg	0.97	0.97	0.97	124
weighted avg	0.97	0.97	0.97	124

Accuracy: 96.7741935483871

```
[55]: LR_roc_auc_score= roc_auc_score(y_test, lr_clf.decision_function(X_test),
    ↪average=None)
print("LR_roc_auc_score :",LR_roc_auc_score )
```

LR\_roc\_auc\_score : 0.9779411764705883

```
[56]: metrics.plot_roc_curve(lr_clf, X_test, y_test)
plt.show()
```



#### 0.0.4 Random Forest Classifier

```
[57]: from sklearn.ensemble import RandomForestClassifier
```

```
rf_clf = RandomForestClassifier( random_state=0).fit(X_train, y_train)
rf_pred = rf_clf.predict(X_test)
rf_yhat_prob = rf_clf.predict_proba(X_test)
```

```
[58]: from sklearn.metrics import classification_report, confusion_matrix,
      ↪ accuracy_score
```

```
print("confusion matrix\n",confusion_matrix(y_test,rf_pred))
print("\nclassification report\n",classification_report(y_test,rf_pred))
print("Accuracy: ",accuracy_score(y_test, rf_pred)*100)
```

confusion matrix

```
[[54  2]
 [ 2 66]]
```

classification report

	precision	recall	f1-score	support
Demented	0.96	0.96	0.96	56


Nondemented	0.97	0.97	0.97	68
accuracy			0.97	124
macro avg	0.97	0.97	0.97	124
weighted avg	0.97	0.97	0.97	124

Accuracy: 96.7741935483871

```
[59]: RF_accuracy_score = accuracy_score(y_test, rf_pred)*100
print("RF_accuracy_score :", RF_accuracy_score )
```

RF\_accuracy\_score : 96.7741935483871

```
[60]: plot_confusion_matrix(rf_clf, X_test, y_test)
plt.show()
```




```
[61]: # Confusion matrix using crosstab method of pandas.
RF_pd = pd.crosstab(y_test, rf_pred, rownames=['True'], colnames=['Predicted'],
↪ margins=True)
RF_pd
```

```
[61]: Predicted   Demented  Nondemented  All
True
Demented         54           2       56
Nondemented      2          66       68
```

```
[62]: RF_roc_auc_score= roc_auc_score(y_test, rf_clf.predict_proba(X_test)[:,-1],  
    ↪average=None)  
print("RF roc_auc score :",RF_roc_auc_score )
```

RF roc\_auc score : 0.9839810924369747

```
[63]: metrics.plot_roc_curve(rf_clf, X_test, y_test)  
  
plt.show()
```



### 0.0.5 Naive Bayes

```
[64]: from sklearn.naive_bayes import GaussianNB
```

```
[65]: gnb_clf = GaussianNB()
```

```
[66]: y_pred_gnb = gnb_clf.fit(X_train, y_train).predict(X_test)
```

```
[67]: #yhat_gnb = gnb.predict(X_test)  
  
gnb_yhat_prob = gnb_clf.predict_proba(X_test)
```

```
[68]: gnb_Accuracy_score = metrics.accuracy_score(y_test, y_pred_gnb)*100
gnb_Accuracy_score
```

```
[68]: 95.96774193548387
```

```
[69]: print("confusion matrix\n",confusion_matrix(y_test,y_pred_gnb))
print("\nclassification report\n",classification_report(y_test,y_pred_gnb))
print("Accuracy: ",accuracy_score(y_test, y_pred_gnb)*100)
```

confusion matrix

```
[[53  3]
```


```
 [ 2 66]]
```

classification report

	precision	recall	f1-score	support
Demented	0.96	0.95	0.95	56
Nondemented	0.96	0.97	0.96	68
accuracy			0.96	124
macro avg	0.96	0.96	0.96	124
weighted avg	0.96	0.96	0.96	124

Accuracy: 95.96774193548387

```
[70]: plot_confusion_matrix(gnb_clf, X_test, y_test)
plt.show()
```

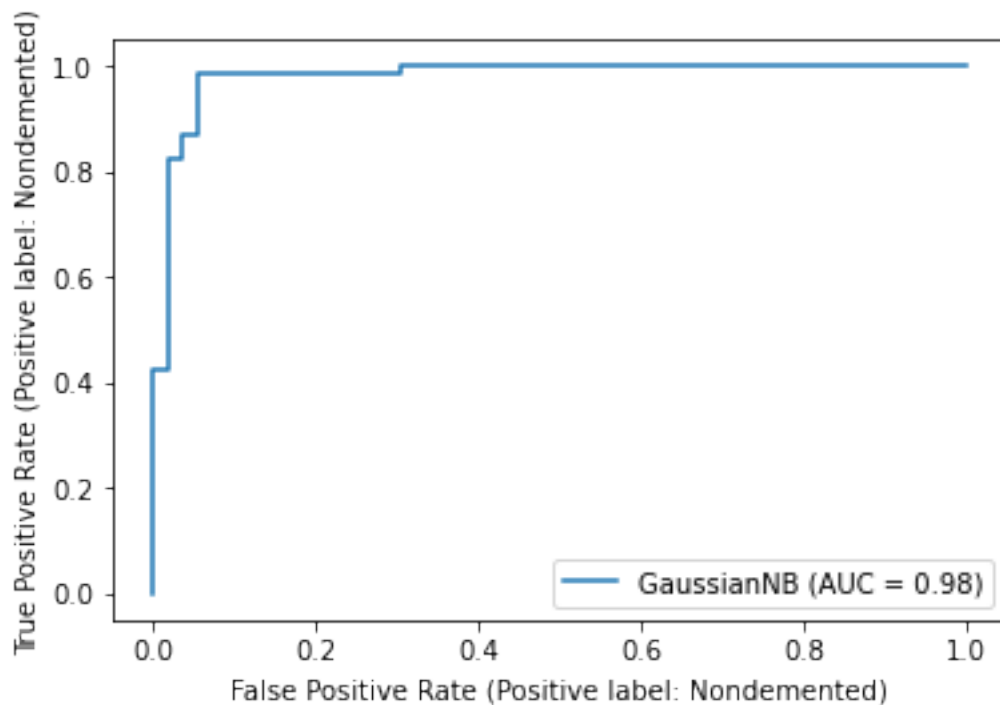


```
[71]: # Confusion matrix using crosstab method of pandas.
GNV_pd = pd.crosstab(y_test, y_pred_gnb, rownames=['True'],
↳ colnames=['Predicted'], margins=True)
GNV_pd
```

```
[71]: Predicted    Demented  Nondemented  All
True
Demented          53           3      56
Nondemented        2          66      68
All                55          69     124
```

```
[72]: metrics.plot_roc_curve(gnb_clf, X_test, y_test)

plt.show()
```



```
[73]: gnb_roc_auc_score= roc_auc_score(y_test, gnb_clf.predict_proba(X_test)[:,-1],
↳ average=None)
print("GNB roc_auc score :", gnb_roc_auc_score )
```

```
GNB roc_auc score : 0.9805672268907563
```

## 0.0.6 Ada Boosting

```
[74]: from sklearn.ensemble import AdaBoostClassifier
```

```
ada_clf = AdaBoostClassifier(n_estimators=100, random_state=42)
ada_clf.fit(X_train, y_train)
```

```
[74]: AdaBoostClassifier(n_estimators=100, random_state=42)
```

```
[75]: yhat_ada = ada_clf.predict(X_test)
```

```
#yhat_ada
```

```
[76]: ada_yhat_prob = ada_clf.predict_proba(X_test)
#ada_yhat_prob
```

```
[77]: ada_accuracy_score = accuracy_score(y_test, yhat_ada)*100
print("ada_accuracy_score :", ada_accuracy_score )
```

```
ada_accuracy_score : 96.7741935483871
```

```
[78]: from sklearn.metrics import classification_report, confusion_matrix, \
      ↪ accuracy_score

print("confusion matrix\n", confusion_matrix(y_test, yhat_ada))
print("\n classification report\n", classification_report(y_test, yhat_ada))
print("Accuracy: ", accuracy_score(y_test, yhat_ada)*100)
```

```
confusion matrix
```

```
[[54  2]
 [ 2 66]]
```


```
classification report
```

	precision	recall	f1-score	support
Demented	0.96	0.96	0.96	56
Nondemented	0.97	0.97	0.97	68
accuracy			0.97	124
macro avg	0.97	0.97	0.97	124
weighted avg	0.97	0.97	0.97	124

```
Accuracy: 96.7741935483871
```

```
[79]: plot_confusion_matrix(ada_clf, X_test, y_test)
plt.show()
```






```
[80]: # Confusion matrix using crosstab method of pandas.
ada_pd = pd.crosstab(y_test, yhat_ada, rownames=['True'],
                    ↳ colnames=['Predicted'], margins=True)
ada_pd
```

```
[80]: Predicted    Demented  Nondemented  All
True
Demented          54           2    56
Nondemented        2          66    68
All                56          68   124
```

```
[81]: metrics.plot_roc_curve(ada_clf, X_test, y_test)

plt.show()
```



```
[82]: ada_roc_auc_score= roc_auc_score(y_test, ada_clf.decision_function(X_test),
    ↪average=None)
    print("ada_roc_auc_score :",ada_roc_auc_score )
```

ada\_roc\_auc\_score : 0.9719012605042017

### 0.0.7 Gradient boosting

```
[83]: from sklearn.ensemble import GradientBoostingClassifier
```

```
[84]: Gra_clf = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0,
    ↪max_depth=1, random_state=0).fit(X_train,
    ↪y_train)
```

```
[85]: Gra_clf.score(X_test, y_test)
```

[85]: 0.9758064516129032

```
[86]: yhat_Gra_cl = Gra_clf.predict(X_test)

    Gra_yhat_prob = Gra_clf.predict_proba(X_test)
```

```
[87]: #print(accuracy_score(y_test, yhat_Gra_cl)*100)
    #print(accuracy_score(y_test, yhat_Gra_cl))
```

```

gra_accuracy_score = Gra_clf.score(X_test, y_test)*100
print("gra_accuracy_score :", gra_accuracy_score )

```

gra\_accuracy\_score : 97.58064516129032

```

[88]: print("confusion matrix\n",confusion_matrix(y_test,yhat_Gra_cl))
      print("\nclassification report\n",classification_report(y_test,yhat_Gra_cl))
      print("Accuracy: ",accuracy_score(y_test, yhat_Gra_cl)*100)

```

confusion matrix

```

[[54  2]
 [ 1 67]]

```

classification report


	precision	recall	f1-score	support
Demented	0.98	0.96	0.97	56
Nondemented	0.97	0.99	0.98	68
accuracy			0.98	124
macro avg	0.98	0.97	0.98	124
weighted avg	0.98	0.98	0.98	124

Accuracy: 97.58064516129032

```

[89]: plot_confusion_matrix(Gra_clf, X_test, y_test)
      plt.show()

```

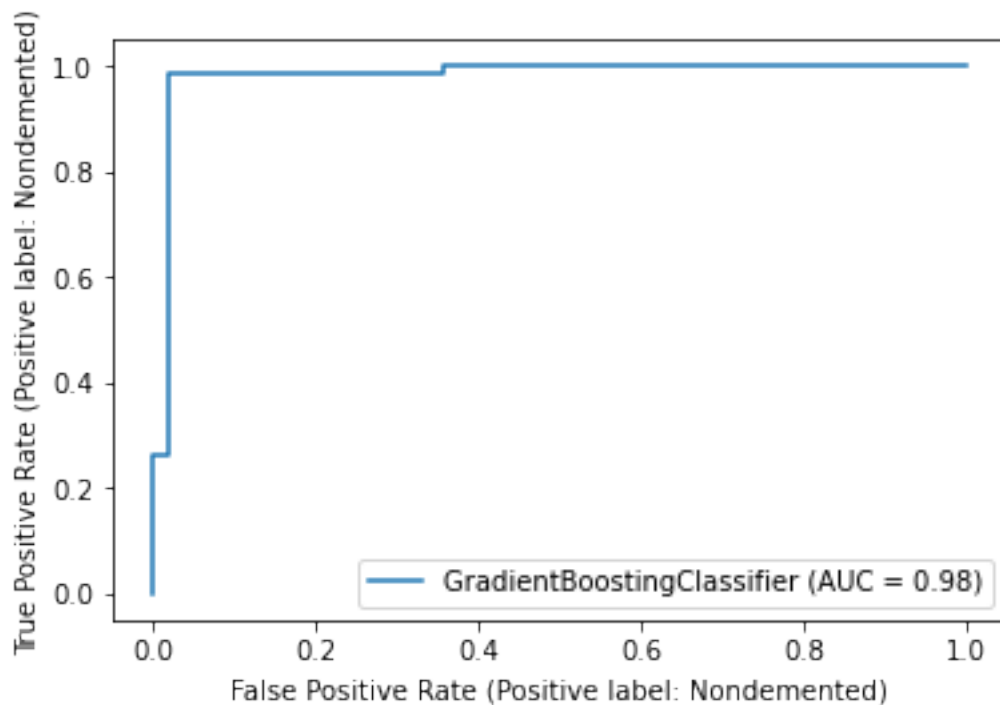


```
[90]: # Confusion matrix using crosstab method of pandas.
Gra_pd = pd.crosstab(y_test, yhat_Gra_cl, rownames=['True'],
    ↳ colnames=['Predicted'], margins=True)
Gra_pd
```

```
[90]: Predicted    Demented  Nondemented  All
True
Demented          54           2    56
Nondemented        1          67    68
All                55          69   124
```

```
[91]: metrics.plot_roc_curve(Gra_clf, X_test, y_test)

plt.show()
```



```
[92]: gra_roc_auc_score= roc_auc_score(y_test, Gra_clf.decision_function(X_test),
    ↳ average=None)
print("gra_roc_auc_score :", gra_roc_auc_score )
```

```
gra_roc_auc_score : 0.9818802521008403
```

## 0.0.8 Results

### Accuracy score report dataframe

```
[102]: # create a accuracy list
accuracy_score = [
    ↪SVM_accuracy_score,LR_accuracy_score,RF_accuracy_score,gnb_Accuracy_score,ada_accuracy_score,gr

# create a AUC-score list
auc_score = [
    ↪[SVM_roc_auc_score,LR_roc_auc_score,RF_roc_auc_score,gnb_roc_auc_score,ada_roc_auc_score,gr

# fomulate the report format
df_report = pd.DataFrame(accuracy_score, index=['SVM','Logistic_
    ↪Regression','Random Forest','Naive Bayes','Ada Boosting','Gradient_
    ↪boosting',])
df_report.columns = ['Accuracy Score']
df_report.insert(loc=1, column='ROC AUC Score', value= auc_score)
df_report.columns.name = 'Algorithm'
df_report
```

```
[102]: Algorithm          Accuracy Score ROC AUC Score
SVM                96.774194      0.968487
Logistic Regression 96.774194      0.977941
Random Forest      96.774194      0.983981
Naive Bayes        95.967742      0.980567
Ada Boosting       96.774194      0.971901
Gradient boosting  97.580645      0.98188
EVM                96.774194      NA
```

```
[103]: fig = metrics.plot_roc_curve(rf_clf, X_test, y_test)
fig = metrics.plot_roc_curve(Gra_clf, X_test, y_test,ax = fig.ax_ )
fig = metrics.plot_roc_curve(ada_clf, X_test, y_test,ax = fig.ax_ )
fig = metrics.plot_roc_curve(gnb_clf, X_test, y_test,ax = fig.ax_ )
#metrics.plot_roc_curve(rf_regressor, X_test, y_test)
fig = metrics.plot_roc_curve(lr_clf, X_test, y_test,ax = fig.ax_ )
fig = metrics.plot_roc_curve(svm_clf, X_test, y_test,ax = fig.ax_ )

fig.figure_.suptitle("ROC curve comparison")
fig.figure_.set_size_inches(14,10)
plt.show()
```

ROC curve comparison

